RILO · Evaluation report

Evaluation: exp_a26e263257194f93

Created: 2026-11-09T18:39:36.482125+00:00 · Status: completed

Task and conditions

tell me about the git hub repo

Category: repository · Model: ollama/qwen2.5-coder:7b

Limits: {"max_model_calls": 2, "max_tool_calls": 2, "timeout_seconds": 180}

Snapshot: snap_c0d31363059c486d · Hash: cd6782445bf531583caf8502b765993738a59e07a4ee783dd33b0588f0274785

Input metadata: {"files_seen": 40, "files_captured": 40, "skipped_files": [], "languages": [".ini", ".json", ".lock", ".md", ".py", ".toml", ".txt", ".yml"], "manifests": ["backend/pyproject.toml", "backend/requirements.txt", "frontend/package.json", "tests/fixtures/vulnerable-python/requirements.txt"], "limitations": ["Analysis covers captured and selected text only.", "Static observations are not proof of a vulnerability.", "Repository code and dependencies are never executed on the host."], "redacted_files": ["backend/app/tests/conftest.py", "backend/app/tests/test_cloud_routing.py", "backend/app/tests/test_integration.py"], "repository": "https://github.com/Gauravwakle96/KYRO", "commit_sha": "a03f97ded1f5da474cf16eefe08afec1b5133e07", "default_branch": "main", "files_in_tree": 319, "retrieval_limit": 40, "omitted_candidate_files": 188, "retrieval": "commit-pinned bounded source files", "static_analysis": {"version": "repository-static-v1", "modules": [{"path": "backend/app/__init__.py", "imports": [], "function_count": 0, "routes": []}, {"path": "backend/app/agents/engine.py", "imports": ["app.agents.models", "app.agents.schemas", "app.core.config", "app.core.security", "app.database", "app.jobs.worker", "app.repository.intelligence", "app.sandbox.runner", "ast", "difflib", "json", "pathlib", "shutil", "sqlalchemy.orm"], "function_count": 7, "routes": []}, {"path": "backend/app/agents/models.py", "imports": ["app.core.config", "app.core.security", "app.database", "httpx", "ipaddress", "json", "pathlib", "socket", "sqlalchemy", "sqlalchemy.orm", "time", "urllib.parse"], "function_count": 4, "routes": []}, {"path": "backend/app/agents/schemas.py", "imports": ["app.sandbox.runner", "pydantic"], "function_count": 1, "routes": []}, {"path": "backend/app/api/routes.py", "imports": ["app.agents.engine", "app.agents.models", "app.api.schemas", "app.core.config", "app.core.security", "app.database", "app.integrations.github", "app.intelligence.assessment", "app.jobs.recovery", "app.sandbox.runner", "app.scanners.adapters", "argon2.exceptions", "asyncio", "datetime", "fastapi", "fastapi.responses", "httpx", "json", "secrets", "shutil", "sqlalchemy", "sqlalchemy.exc", "sqlalchemy.orm", "threading"], "function_count": 61, "routes": [{"function": "health", "line": 126, "method": "get", "path": "/health"}, {"function": "setup", "line": 132, "method": "get", "path": "/auth/setup"}, {"function": "register", "line": 137, "method": "post", "path": "/auth/register"}, {"function": "login", "line": 149, "method": "post", "path": "/auth/login"}, {"function": "me", "line": 164, "method": "get", "path": "/auth/me"}, {"function": "logout", "line": 169, "method": "post", "path": "/auth/logout"}, {"function": "projects", "line": 177, "method": "get", "path": "/projects"}, {"function": "create_project", "line": 182, "method": "post", "path": "/projects"}, {"function": "update_project", "line": 191, "method": "patch", "path": "/projects/{project_id}"}, {"function": "dele

te_project", "line": 200, "method": "delete", "path": "/projects/{project_id}"}, {"function": "github_status", "line": 224, "method": "get", "path": "/github/status"}, {"function": "github_device", "line": 230, "method": "post", "path": "/github/device"}, {"function": "github_poll", "line": 244, "method": "post", "path": "/github/device/poll"}, {"function": "github_disconnect", "line": 267, "method": "delete", "path": "/github/connection"}, {"function": "github_repos", "line": 277, "method": "get", "path": "/github/repositories"}, {"function": "github_branches", "line": 286, "method": "get", "path": "/github/branches"}, {"function": "repositories", "line": 297, "method": "get", "path": "/repositories"}, {"function":

"connect_repository", "line": 302, "method": "post", "path": "/repositories"}, {"function": "fixture", "line": 311, "method": "post", "path": "/repositories/fixture"}, {"function": "repository", "line": 320, "method": "get", "path": "/repositories/{repository_id}"}, {"function": "analyze", "line": 325, "method": "post", "path": "/repositories/{repository_id}/analyze"}, {"function": "source_file", "line": 333, "method": "get", "path": "/repositories/{repository_id}/files"}, {"function": "repository_symbols", "line": 342, "method": "get", "path": "/repository-intelligence/{repository_id}"}, {"function": "scan", "line": 348, "method": "post", "path": "/repositories/{repository_id}/scans"}, {"function": "scanners", "line": 356, "method": "get", "path": "/scanners"}, {"function": "scans", "line": 361, "method": "get", "path": "/scans"}, {"function": "scan_details", "line": 366, "method": "get", "path": "/scans/{scan_id}"}, {"function": "findings", "line": 375, "method": "get", "path": "/findings"}, {"function": "finding_details", "line": 388, "method": "get", "path": "/findings/{finding_id}"}, {"function": "enrich_finding", "line": 393, "method": "post", "path": "/findings/{finding_id}/enrich"}, {"function": "dependency_list", "line": 399, "method": "get", "path": "/dependencies"}, {"function": "graph", "line": 404, "method": "get", "path": "/attack-paths/{repository_id}"}, {"function": "remediation_plan", "line": 410, "method": "post", "path": "/findings/{finding_id}/remediation"}, {"function": "remediation_list", "line": 419, "method": "get", "path": "/remediation"}, {"function": "remediation_details", "line": 424, "method": "get", "path": "/remediation/{task_id}"}, {"function": "approve", "line": 430, "method": "post", "path": "/remediation/{task_id}/approve"}, {"function": "execute", "line": 442, "method": "post", "path": "/remediation/{task_id}/execute"}, {"function": "verification", "line": 452, "method": "post", "path": "/remediation/{task_id}/verify"}, {"function": "report", "line": 460, "method": "get", "path": "/remediation/{task_id}/report"}, {"function": "deliver", "line": 468, "method": "post", "path": "/remediation/{task_id}/pull-request"}, {"function": "pull_requests", "line": 480, "method": "get", "path": "/pull-requests"},

{"function": "refresh_pr", "line": 485, "method": "post", "path": "/remediation/{task_id}/pull-request/refresh"}, {"function": "models", "line": 501, "method": "get", "path": "/models"}, {"function": "model_settings", "line": 506, "method": "put", "path": "/models/settings"}, {"function": "jobs", "line": 518, "method": "get", "path": "/jobs"}, {"function": "job_details", "line": 523, "method": "get", "path": "/jobs/{job_id}"}, {"function": "cancel_job", "line": 531, "method": "post", "path": "/jobs/{job_id}/cancel"}, {"function": "retry_job", "line": 546, "method": "post", "path": "/jobs/{job_id}/retry"}, {"function": "job_events", "line": 575, "method": "get", "path": "/jobs/{job_id}/events"}, {"function": "analytics", "line": 601, "method": "get", "path": "/analytics"}, {"function": "configuration", "line": 613, "method": "get", "path": "/settings"}, {"function": "audit_history", "line": 618, "method": "get", "path": "/audit"}]}, {"path": "backend/app/api/schemas.py", "imports": ["app.integrations.github", "pydantic", "re", "typing"], "function_count": 2, "routes": []}, {"path": "backend/app/core/config.py", "imports": ["pathlib", "pydantic", "pydantic_settings", "typing"], "function_count": 1, "routes": []}, {"path": "backend/app/core/errors.py", "imports": [], "function_count": 1, "routes": []}, {"path": "backend/app/core/security.py", "imports": ["app.core.config", "app.database", "argon2", "cryptography.fernet", "datetime", "fastapi", "hashlib", "os", "pathlib", "re", "secrets", "sqlalchemy.orm"], "function_count": 5, "routes": []}, {"path": "backend/app/database.py", "imports": ["app.core.config", "collections.abc", "datetime", "sqlalchemy", "sqlalchemy.orm", "uuid"], "function_count": 5, "routes": []}, {"path": "backend/app/integrations/github.py", "imports": ["app.core.config", "app.core.errors", "app.core.security", "app.database", "base64", "httpx", "io", "jwt", "pathlib", "re", "sqlalchemy.orm", "tarfile", "time", "urllib.parse"], "function_count": 8, "routes": []}, {"path": "backend/app/intelligence/assessment.py", "imports": ["app.database", "sqlalchemy", "sqlalchemy.orm"], "function_count": 3, "routes": []}, {"path": "backend/app/intelligence/enrichment.py", "imports": ["app.core.config", "app.core.security", "app.database", "app.jobs.recovery", "httpx", "json", "re", "time", "urllib.parse"], "function_count": 2, "routes": []}, {"path": "backend/app/jobs/recovery.py", "imports": ["app.core.config", "app.core.errors", "app.core.security", "app.database", "datetime", "httpx"], "function_count": 5, "routes": []}, {"path": "backend/app/jobs/worker.py", "imports": ["app.core.config", "app.core.security", "app.database", "app.jobs", "collections.abc", "datetime", "sqlalchemy", "threading"], "function_count": 8, "routes": []}, {"path": "backend/app/main.py", "imports": ["alembic", "alembic.config", "app.agents.models", "app.api.routes", "app.core.config", "app.core.security", "app.jobs", "app.services.handlers", "collections", "contextli

b", "fastapi", "fastapi.middleware.cors", "fastapi.middleware.trustedhost", "fastapi.responses", "fastapi.staticfiles", "json", "logging", "pathlib", "shutil", "threading", "time"], "function_count": 7, "routes": [{"function": "api_documentation", "line": 120, "method": "get", "path": "/docs"}, {"function": "spa", "line": 130, "method": "get", "path": "/{path:path}"}]}, {"path": "backend/app/repository/intelligence.py", "imports": ["app.core.security", "app.database", "ast", "collections", "json", "packaging.requirements", "pathlib", "re", "sqlalchemy", "sqlalchemy.orm", "tomllib", "tree_sitter", "tree_sitter_javascript", "tree_sitter_python", "tree_sitter_typescript"], "function_count": 5, "routes": []}, {"path": "backend/app/sandbox/runner.py",

"imports": ["app.core.config", "app.core.security", "app.jobs.worker", "json", "pathlib", "re", "subprocess", "tempfile", "time", "uuid"], "function_count": 9, "routes": []}, {"path": "backend/app/scanners/adapters.py", "imports": ["app.core.config", "app.sandbox.runner", "app.scanners.normalization", "json", "pathlib", "uuid"], "function_count": 3, "routes": []}, {"path": "backend/app/scanners/normalization.py", "imports": ["app.core.security", "re"], "function_count": 7, "routes": []}, {"path": "backend/app/scanners/service.py", "imports": ["app.database", "app.intelligence.assessment", "app.jobs.worker", "app.scanners.adapters", "app.scanners.normalization", "pathlib", "sqlalchemy", "sqlalchemy.orm"], "function_count": 1, "routes": []}, {"path": "backend/app/services/handlers.py", "imports": ["app.agents.engine", "app.core.errors", "app.core.security", "app.database", "app.integrations.github", "app.intelligence.assessment", "app.intelligence.enrichment", "app.jobs.worker", "app.repository.intelligence", "app.scanners.service", "app.verification.engine", "pathlib", "shutil", "sqlalchemy.orm"], "function_count": 5, "routes": []}, {"path": "backend/app/tests/conftest.py", "imports": ["app.database", "app.main", "fastapi.testclient", "os", "pytest", "sqlalchemy", "tempfile"], "function_count": 4, "routes": []}, {"path": "backend/app/tests/test_agent_reliability.py", "imports": ["app.agents", "app.agents.engine", "app.agents.schemas", "app.core.config", "app.core.security", "app.database", "httpx", "json", "pytest", "sqlalchemy", "sqlalchemy.exc"], "function_count": 8, "routes": []}, {"path": "backend/app/tests/test_analytics_counts.py", "imports": ["app.database"], "function_count": 1, "routes": []}, {"path": "backend/app/tests/test_api_reference.py", "imports": [], "function_count": 1, "routes": []}, {"path": "backend/app/tests/test_automatic_recovery.py", "imports": ["app.core.config", "app.core.errors", "app.database", "app.integrations.github", "app.jobs", "app.services", "datetime", "httpx", "pytest"], "function_count": 21, "routes": []}, {"path": "backend/app/tests/test_cloud_routing.py", "imports": ["app.agents", "app.core.config", "app.database", "httpx", "json", "pytest", "sqlalchemy"], "functio

n_count": 3, "routes": []}, {"path": "backend/app/tests/test_contracts.py", "imports": ["app.agents", "app.core.config", "app.database", "app.integrations.github", "app.sandbox.runner", "app.scanners.adapters", "app.scanners.normalization", "httpx", "io", "json", "pytest", "sqlalchemy", "sys", "tarfile"], "function_count": 22, "routes": []}, {"path": "backend/app/tests/test_integration.py", "imports": ["app.agents.models", "app.core.config", "app.core.security", "app.database", "app.jobs", "app.repository.intelligence", "pytest"], "function_count": 9, "routes": []}, {"path": "backend/main.py", "imports": ["app.main"], "function_count": 0, "routes": []}, {"path": "backend/migrations/env.py", "imports": ["alembic", "app.database"], "function_count": 0, "routes": []}], "observations": [], "dependencies": [{"requirement": "fastapi>=0.121", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "uvicorn[standard]>=0.38", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "sqlalchemy>=2.0.44", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "alembic>=1.17", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "pydantic-settings>=2.11", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "httpx>=0.28", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "argon2-cffi>=25.1", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "cryptography>=46", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "pyjwt>=2.10", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "psycopg[binary]>=3.2", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "tree-sitter==0.25.2", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "tree-sitter-python>=0.25", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "tree-sitter-javascript>=0.25", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "tree-sitter-typescript>=0.23", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"requirement": "packaging>=25", "ecosystem": "PyPI", "manifest": "backend/pyproject.toml", "vulnerability_status": "not checked"}, {"name": "@fontsource-variable/inter", "version_spec": "^5.2.8", "ecosystem": "npm", "manifest": "frontend/package.json", "

vulnerability_status": "not checked"}, {"name": "@fontsource/jetbrains-mono", "version_spec": "^5.2.8", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@tanstack/react-query", "version_spec": "^5.90.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "lucide-react", "version_spec": "^0.577.0", "ecosystem": "npm", "manifest": "frontend/package.json",

"vulnerability_status": "not checked"}, {"name": "react", "version_spec": "^19.2.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "react-dom", "version_spec": "^19.2.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "react-router-dom", "version_spec": "^7.9.6", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@axe-core/playwright", "version_spec": "^4.11.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@eslint/js", "version_spec": "^9.39.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@playwright/test", "version_spec": "^1.56.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@testing-library/jest-dom", "version_spec": "^6.9.1", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@testing-library/react", "version_spec": "^16.3.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@types/react", "version_spec": "^19.2.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@types/react-dom", "version_spec": "^19.2.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "@vitejs/plugin-react", "version_spec": "^5.1.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "eslint", "version_spec": "^9.39.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "eslint-plugin-react-hooks", "version_spec": "^7.0.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "eslint-plugin-react-refresh", "version_spec": "^0.4.24", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "globals", "version_spec": "^16.5.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "jsdom", "version_spec": "^27.2.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "typescript", "version_spec": "^5.9.3", "ecosystem":

"npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "typescript-eslint", "version_spec": "^8.46.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "vite", "version_spec": "^7.2.2", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}, {"name": "vitest", "version_spec": "^4.0.0", "ecosystem": "npm", "manifest": "frontend/package.json", "vulnerability_status": "not checked"}], "test_files": ["backend/app/tests/conftest.py", "backend/app/tests/test_agent_reliability.py", "backend/app/tests/test_analytics_counts.py", "backend/app/tests/test_api_reference.py", "backend/app/tests/test_automatic_recovery.py", "backend/app/tests/test_cloud_routing.py", "backend/app/tests/test_contracts.py", "backend/app/tests/test_integration.py"], "executed_repository_code": false, "limitations": ["Static rule matches require contextual review.", "Dependency versions are declarations, not installed or resolved versions.", "Known vulnerability and outdated-package lookups have not been performed.", "Only captured files were inspected."]}, "runtime_tests": {"status": "failed", "executed": true, "exit_code": 5, "stdout": "", "stderr": "\n--\nRan 0 tests in 0.000s\n\nNO TESTS RAN\n", "command": ["python", "-I", "-m", "unittest", "discover", "-v"], "image": "python@sha256:e48f59cdba69a3cbd9ed2cf0f348d90d45a8d350d98134cf1070bc106ee9ed90", "network": "none", "timeout_seconds": 60, "scope": "Captured source subset only; no dependency installation", "cleanup": {"status": "completed", "container": "rilo-test-79b86f2ec828", "reason": null}}}

Evaluation rubric: {"version": "workspace-evidence-v1", "category": "repository", "ground_truth": null, "accuracy": "undefined without labels"}

Execution order: sequential, independent agent histories. Shared captured evidence.

Relay

Run: run_2de1402dfda240e4 · completed

Latency: 38439.05 ms · Input tokens: 932 · Output tokens: 482 · API cost: 0.0

Null usage or cost means unknown/not applicable; consult provenance in JSON.

The KYRO GitHub repository is described in the README.md file. It is an AI-powered autonomous software security and engineering intelligence tool designed for local execution and human-reviewed changes. The tool follows a workflow from Understand to Deliver, and it requires approval before editing a repository workspace and publishing a draft pull request.

Task checks · Promptfoo 0.125.1

Status: ungraded · Score: None · Explicit task checks only; not a measurement of overall factual accuracy.

finding_1: Project Description [info]

KYRO is an AI-powered autonomous software security and engineering intelligence tool.

This information is directly stated in the README.md file.

PARTIALLY VERIFIED · Reference integrity only; interpretation needs review · Method: snapshot_reference_and_exact_passage

Source: README.md · Lines 3–5 · Page None · Valid reference: True

Think. Build. Secure.

AI-powered autonomous software security and engineering intelligence, built for local execution and human-reviewed changes.

Matching method: exact passage

finding_2: Workflow [info]

KYRO follows a workflow from Understand to Deliver.

This information is explicitly mentioned in the README.md file.

PARTIALLY VERIFIED · Reference integrity only; interpretation needs review · Method: snapshot_reference_and_exact_passage

Source: README.md · Lines 9–9 · Page None · Valid reference: True

Its workflow is **Understand → Detect → Investigate → Prioritize → Plan → Approve → Remediate → Test → Verify → Deliver**.

Matching method: exact passage

finding_3: Approval Requirement [info]

Approval is required before editing a repository workspace and publishing a draft pull request.

This requirement is clearly stated in the README.md file.

PARTIALLY VERIFIED · Reference integrity only; interpretation needs review · Method: snapshot_reference_and_exact_passage

Source: README.md · Lines 9–9 · Page None · Valid reference: True

Approval is required before editing a repository workspace and before publishing a draft pull request.

Matching method: exact passage

Metrics: {"citation_validity": 1.0, "evidence_coverage": 1.0, "precision": null, "recall": null, "f1": null, "accuracy": null}

Limitation: No ground-truth labels: accuracy, precision, recall and F1 are undefined.

Limitation: Valid references do not prove an interpretation. Free-text summary is not fully verified.

Limitation: Subjective quality and instruction adherence need independent review.

Repo Analyst

Run: run_a1d8fe8306224902 · completed

Latency: 65348.24 ms · Input tokens: 1963 · Output tokens: 900 · API cost: 0.0

Null usage or cost means unknown/not applicable; consult provenance in JSON.

The KYRO repository is an AI-powered autonomous software security and engineering intelligence tool designed for local execution and human-reviewed changes. It follows a workflow from Understand to Deliver, involving various stages like Detect, Investigate, Prioritize, Plan, Approve, Remediate, Test, and Verify. The repository includes a README.md file that outlines the project's purpose, capabilities, and workflow. The product is not certified secure or ready for public multi-tenant operation.

Task checks · Promptfoo 0.125.1

Status: ungraded · Score: None · Explicit task checks only; not a measurement of overall factual accuracy.

finding_1: Project Purpose [info]

The project aims to provide AI-powered autonomous software security and engineering intelligence.

This information is stated directly in the README.md file and is crucial for understanding the project's objectives.

PARTIALLY VERIFIED · Reference integrity only; interpretation needs review · Method: snapshot_reference_and_exact_passage

Source: README.md · Lines 3–5 · Page None · Valid reference: True

Think. Build. Secure.

AI-powered autonomous software security and engineering intelligence, built for local execution and human-reviewed changes.

Matching method: exact passage

finding_2: Workflow Stages [info]

The KYRO workflow includes stages such as Understand, Detect, Investigate, Prioritize, Plan, Approve, Remediate, Test, and Verify.

This information is explicitly mentioned in the README.md file and is essential for understanding the project's operational flow.

PARTIALLY VERIFIED · Reference integrity only; interpretation needs review · Method: snapshot_reference_and_exact_passage

Source: README.md · Lines 9–9 · Page None · Valid reference: True

Its workflow is **Understand → Detect → Investigate → Prioritize → Plan → Approve → Remediate → Test → Verify → Deliver**.

Matching method: exact passage

finding_3: Product Status [info]

The product is not certified secure or ready for public multi-tenant operation.

This statement is made in the README.md file and is important for users to be aware of the product's current state.

PARTIALLY VERIFIED · Reference integrity only; interpretation needs review · Method: snapshot_reference_and_exact_passage

Source: README.md · Lines 13–13 · Page None · Valid reference: True

This product is not certified secure or ready for public multi-tenant operation.

Matching method: exact passage

Metrics: {"citation_validity": 1.0, "evidence_coverage": 1.0, "precision": null, "recall": null, "f1": null, "accuracy": null}

Limitation: No ground-truth labels: accuracy, precision, recall and F1 are undefined.

Limitation: Valid references do not prove an interpretation. Free-text summary is not fully verified.

Limitation: Subjective quality and instruction adherence need independent review.

Comparison and retest

{"method": "promptfoo-explicit-checks-v1", "scope": "Explicit task checks only; not a measurement of overall factual accuracy.", "agents": [{"run_id": "run_2de1402dfda240e4", "agent_id": "relay", "agent_name": "Relay", "outcome": "completed", "status": "ungraded", "score": null, "passed": null, "total": null, "duration_ms": 38439.05, "reason": null}, {"run_id": "run_a1d8fe8306224902", "agent_id": "repo_analyst", "agent_name": "Repo Analyst", "outcome": "completed", "status": "ungraded", "score": null, "passed": null, "total": null, "duration_ms": 65348.24, "reason": null}], "winner": null, "verdict": "ungraded", "explanation": "No success checks were defined. Compare answers, evidence and usage; accuracy and a winner remain undetermined. Add explicit checks to the next task.", "findings": [{"agent_a": "relay", "agent_b": "repo_analyst", "finding_a": "finding_2", "finding_b": "finding_2", "relationship": "potential disagreement", "outcome": "Inconclusive", "method": "shared_source_and_title_overlap", "explanation": "Related findings share source references. Different wording alone does not establish a contradiction or a winner."}]}

Retest against the same snapshot, prompt, model, budgets and rubric. Review changed conditions before interpreting a difference. Source integrity does not establish claim correctness. Unverified findings need review.